



ELSEVIER

Available online at [www.sciencedirect.com](http://www.sciencedirect.com)

**SciVerse ScienceDirect**

**Electronic Notes in  
Theoretical Computer  
Science**

Electronic Notes in Theoretical Computer Science 279 (3) (2011) 73–84

[www.elsevier.com/locate/entcs](http://www.elsevier.com/locate/entcs)

# Generative Version of the FastFlow Multicore Library

Zalán Szűgyi<sup>1</sup> Norbert Pataki<sup>2</sup>

*Department of Programming Languages and Compilers  
Eötvös Loránd University  
Budapest, Hungary*

---

## Abstract

Nowadays, one of the most important challenges in programming is the efficient usage of multicore processors. Many new programming languages and libraries support multicore programming. FastFlow is one of the most promising multicore C++ libraries. Unfortunately, a design problem occurs in the library. One of the most important methods is pure virtual function in a base class. This method supports the communication between different threads. Although, it cannot be template function because of the virtuality, hence, the threads pass and take argument as a `void*` pointer. The base class is not template neither. This is not typesafe approach. We make the library more efficient and safer with the help of generative technologies.

*Keywords:* multicore programming, C++, FastFlow, template

---

## 1 Introduction

The recent trend to increase core count in processors has led to a renewed interest in the design of both methodologies and mechanisms for the effective parallel programming of shared memory computer architectures. Those methodologies are largely based on traditional approaches of parallel computing.

Usually, low-level approaches supplies the programmers only with primitives for flows-of-control management (creation, destruction), their synchronization and data sharing, which are usually accomplished in critical regions accessed in mutual exclusion (mutex). For instance, POSIX thread library can be used to this purpose. Programming parallel complex applications in this way is certainly hard; tuning them for performance is often even harder due to the non-trivial effects induced by memory fences (used to implement mutex) on data replicated in the core's caches.

---

<sup>1</sup> Email: [lupin@ludens.elte.hu](mailto:lupin@ludens.elte.hu)

<sup>2</sup> Email: [patakino@elte.hu](mailto:patakino@elte.hu)

Indeed, memory fences are one of the key sources of performance degradation in communication intensive (e.g. streaming) parallel applications. Avoiding memory fences means not only avoiding locks but also avoiding any kind of atomic operation in memory (e.g. Compare-And-Swap, Fetch-and-Add). While there exists several assessed fence-free solutions for asynchronous symmetric communications, these results cannot be easily extended to asynchronous asymmetric communications that are necessary to support arbitrary streaming networks.

The important approach to ease programmer's task and improve program efficiency consist in to raise the level of abstraction of concurrency management primitives. For example, threads might be abstracted out in higher-level entities that can be pooled and scheduled in user space possibly according to specific strategies to minimize cache flushing or maximize load balancing of cores. Synchronization primitives can be also abstracted out and associated to semantically meaningful points of the code, such as function calls and returns, loops, etc.

This kind of abstraction significantly simplify the hand-coding of applications. However, it is still too low-level to effectively automatize the optimization of the parallel code: the most important weakness here is in the lack of information concerning the intent of the code (idiom recognition); inter-procedural/component optimization further exacerbates the problem.

Recently, there has been a trend of generating programs from high-level specifications. This is called the generative approach, which focuses on synthesizing implementations from higher-level specifications rather than transforming them. From this approach, programmers' goal is captured by the specification. In addition, technologies for code generation are well developed (staging, partial evaluation, automatic programming, generative programming) [7]. FastFlow [4], TBB [13] and OpenMP [8] follow this approach. The programmer needs to explicitly define parallel behaviour by using proper constructs, which clearly bound the interactions among flows-of-control, the read-only data, the associativity of accumulation operations and the concurrent access to shared data structures.

FastFlow is a parallel programming framework for multi-core platforms based upon non-blocking lock-free/fence-free synchronization mechanisms. The framework is composed of a stack of layers that progressively abstracts out the programming of shared-memory parallel applications. The stack has two different goals: to ease the development of applications and make them very fast and scalable. FastFlow is particularly targeted to the development of streaming applications.

*Templates* are key elements of C++ programming language [15]. They enable data structures and algorithms be parameterized by types thus capturing commonalities of abstractions at compile time without performance penalties at runtime. *Generic programming*, a recently emerged programming paradigm for writing reusable components – most cases data structures and algorithms – is implemented in C++ with heavy use of templates.

The worker threads of FastFlow are passing data between each other via `void*` pointers, which are then processed by the receiving thread. This is rather necessary for implementation reasons. At the core of the FastFlow framework there is an

abstract class `ff_node`. It has a pure virtual method `svc`, which is the main method of every thread in the application. However, it is needed to handle any kind of data type, but while it is a virtual function it cannot be template. Although, the base class could be a class template. However, this solution is neither elegant nor type-safe [17].

In this paper we present our solution about how FastFlow library can be extended to be type-safe. We make the abstract base class be templated and we use static polymorphism instead of dynamic polymorphism by curiously recurring template pattern. This way in one hand the system will be type-safe, and on the other hand the virtuality of the method can be eliminated, thus the system becomes more effective.

This paper is organized as follows. In section 2 the FastFlow library is shown and present typical examples that are risky from the view of program correctness. We present an approach that can be used to overcome this problem by the type system of C++ in section 3, the implementation details are described in section 4. The sections 5 and 6 present our pipeline and buffer implementation. In chapter 7 we demonstrate the usage of our extended FastFlow library through an example. Finally, this paper concludes in section 8.

## 2 FastFlow

FastFlow is one of the most efficient multicore C++ frameworks. It is similar to Intel's TBB, advocates a pattern-based (skeletal) approach to parallel programming. FastFlow aims to provide a set of low-level mechanisms able to support low latency and high-bandwidth data flows in a network of threads running on a SCM. These flows, as typical in streaming applications, are supposed to be mostly unidirectional and asynchronous. On these architectures, the key issues regard memory fences, which are required to keep the various caches coherent. The library provides the programmer with two basic mechanisms: Multiple-Producer-Multiple-Consumer (MPMC) queues and a memory allocator. The memory allocator is build on top of MPMC queues and can be substituted either with OS standard allocator (paying a performance penalty) or a third-party allocator (e.g. Intel TBB scalable allocator).

The key intuition underneath FastFlow is to provide the programmer with lock-free MP queues and MC queues (that can be used in pipeline to build MPMC queues) to support fast streaming networks. Traditionally, MPMC queues are build as passive entities: threads concurrently synchronize to access data. Synchronizations are usually supported by one or more atomic operations (e.g. Compare-And-Swap) that behave as memory fences. FastFlow design follows a different approach: the synchronizations among queue readers or writers are arbitrated by an active entity (e.g. a thread). We call these entities Emitter (E) or Collector (C) according to their role; they actually read an item from one or more lock-free Simple-Producer-Simple-Consumer (SPSC) queues and write onto one or more lock-free SPSC queues. Hence, this requires a memory copy. On the other hand, no atomic operation is needed.

The performance advantage of this solution descend from the higher speed of the copy with respect to the memory fence, that advantage is further increased by avoiding cache invalidation triggered by fences. This also depends on the size and the memory layout of copied data. The former point is addressed using data pointers instead of data, and enforcing that the data is not concurrently written: in many cases this can be derived by the semantics of the skeleton that has been implemented using MPMC queues (as an example this is guaranteed in a stateless farm and many other cases).

When using dynamically allocated memory, the memory allocator plays an important role in term of performance. Dynamic memory allocators (`malloc/ free`) rely on mutual exclusion locks for protecting the consistency of their shared data structures under multi-threading. Therefore, the use of memory allocator may subtly reintroduce the locks in the lock-free application. FastFlow includes its own custom memory allocator, which is specifically optimized for SPMC pattern. The basic assumption is that, in streaming application, typically, one thread allocate memory and one or many other threads free memory. This assumption permits to develop a multi-threaded memory allocator that use SPSC channels between the allocator thread and the generic thread that performs the free, avoiding the use of costly lock based protocols for maintaining the memory consistency of the internal structures. However, the FastFlow allocator is not a general purpose allocator and it has limitations.

Let us consider the following code snippet that takes advantage of FastFlow library: The `Emitter` feeds the workers with data, now it just send integers. The `Worker` processes the data, and passes it toward to the other `Worker` or a `Collector`, depends on the initialization of FastFlow. Now it just prints the received integer value to the screen, and pass it toward to the `Collector`. The `Collector` gathers the computation results done by the Workers. Now it just receives the integer value and does nothing.

```
// the generic worker
struct Worker: ff::ff_node
{
    void * svc( void * task )
    {
        int * t = (int *)task;
        std::cout << "Worker " << ff_node::get_my_id()
                  << " received task " << *t << "\n";
        return task;
    }
};

// the gatherer filter
struct Collector: ff::ff_node
{
    void * svc( void * task )
    {
        int * t = (int *)task;
        if (*t == -1) return NULL;
        return task;
    }
};

// the load-balancer filter
struct Emitter: ff::ff_node
{
    Emitter( int max_task ):ntask( max_task ) {}
```

```

void * svc( void * )
{
    int * task = new int( ntask );
    --ntask;
    if ( ntask < 0 ) return NULL;
    return task;
}
private:
    int ntask;
};

int main()
{
    ff_farm<> farm;

    Emitter E( streamlen );
    farm.add_emitter( &E );

    std::vector<ff_node *> w;
    for( int i = 0; i < nworkers; ++i )
        w.push_back( new Worker );
    farm.add_workers( w );

    Collector C;
    farm.add_collector( &C );

    farm.run_and_wait_end();
}

```

The communication between the emitter, the collector and the workers is done via `void*` pointers, so arbitrary data can be passed. But this approach is not typesafe at all. FastFlow works this way because the `svc` method is specified as a pure virtual function in the base class. A template method can be typesafe and elegant solution to pass arbitrary data, however, a (pure) virtual method cannot be template. On the other hand, the abstract base class itself could be template, but the original implementation does not take advantage of this feature.

### 3 Base of the Implementation

Curiously recurring template pattern can be used for static polymorphism [5]. This idiom is able to simulate dynamic binding without virtual functions [9]. In this case, the method is selected at compilation time, hence this approach is more effective at run-time, too. This pattern can be used in many ways [12]. Typical appearance of this approach is the following:

```

template <class Derived>
struct Base
{
    void f()
    {
        static_cast<Derived*>( this )->f();
    }
};

struct Derived: Base<Derived>
{
    void f()
    {
        // actual implementation
    }
};

```

Certainly, this pattern has limitation: the derived types cannot depend on run-time information. However, in our case, we can distinguish between emitter, collector and worker types at compilation time, thus we do not need to deal with runtime information. Therefore, this limitation has no effect on FastFlow's usage.

## 4 Implementation Details

In FastFlow the base class of parallel computing is `ff_node`. That class has a pure virtual method called `svc`, which has a `void*` argument. All of the classes take part in parallel computation must be derived from `ff_node` and must implement the member function `svc`. As the argument type of `svc` is `void*`, it can accept arbitrary data, however, it is not a type safe solution. Inside the function we need to cast the argument to a proper type. Casting to an improper type causes a runtime error. Nevertheless it can be compiled without any error or warning message, making the bugfix more difficult.

Our solution based on curiously recurring template pattern described in section 3. This way, the function `svc` does not need to be virtual and the type of its argument can be a properly typed pointer. The type of the argument is a template parameter of class `ff_node`. The class `ff_node` has another template parameter which refers to the type of its derived class. The following code snippet shows the declaration of class `ff_node`:

```
template<class Derived, class T>
class ff_node;
```

There are three types of nodes in FastFlow: the `emitter`, the `worker` and the `collector` or `fallback`. The tasks are scheduled by the emitter. The workers compute them parallelly and independently. Finally, the collector gathers the results.

The class `ff_farm` represents a farm skeleton. Originally that class has two template arguments, the type of load balancer and the type of gatherer. These template arguments have default value, thus the programmer can skip them and let the FastFlow use the default types. We defined five template arguments of `ff_farm` additionally. The first three template arguments refer to the type of `emitter`, `worker` and `collector`, the fourth one refers the type passed as argument of member function `svc`, the fifth one is a boolean value which sets whether FastFlow uses circular or non-circular buffer. The last two are the original template arguments. The details of buffers are described in section 6. As the `ff_farm` is a child class of `ff_node`, it needs to pass itself to a template argument of `ff_node`. The following code snippet shows the declaration of the class `ff_farm`.

```
template<class E,
        class W,
        class C,
        class T,
        bool is_circular = false,
        class lb=ff_loadbalancer<E, W, C, T>,
        class gt=ff_gatherer<C, W, T> >
class ff_farm:
    public ff_node<ff_farm<E, W, C, T, is_circular, lb, gt>, T>;
```

There are several farm schemes exist which do not require emitter or collector types. It is not problem in the original implementation of FastFlow. The programmer just skips to add emitter or collector to the farm. With our solution the programmer needs to specify the type of all kinds of nodes: emitter, worker and collector. To overcome this problem we created a class called `Void`. This class is also a child of class `ff_node`, and it has a template parameter which denotes the type of the argument of member function `svc`. This class implements the `svc` mem-

ber function but in this implementation that function does nothing. Instantiating `ff_farm`, the programmer can use `Void` type to indicate there is no need for emitter or collector types. The definition of class `Void` is below:

```
template<class T>
struct Void : ff::ff_node<Void, T>
{
    T* svc( T* task )
    {
        return 0;
    }
};
```

## 5 Pipelines

The high level pipeline schema is also supported by FastFlow. The frame class of pipeline is called `ff_pipeline`. The pipeline contains stages, one stage can be either a standalone `ff_node` or a farm skeleton. In the original version, the pipeline stores the stages in a vector of `ff_node*`, where the `ff_node` is the base type of the stages, and the virtual function `svc` does the computation.

In our solution, where the dynamic polymorphism replaced with static polymorphism, this approach does not work. While the stage classes have to add themselves as a template argument of `ff_node`, their base types will be different. Hereby they cannot be stored in one vector.

To overcome this problem we defined the class `ff_pipeline` as template class, and its template argument, is a `boost::mpl::list` [1,20]. This template argument defines the type of the stages. The vector stores `void*` pointers and a template metaprogram – with the help of template argument – will cast them to the proper type on usage. This casting process will be done at compilation time, keeping the type safe, and it is hidden from the programmer.

In the original version of FastFlow the class `ff_pipeline` deals with stages by the interface of `ff_node` and virtual functions. See the example below:

```
for( int i = 0; i < number_of_stages; ++i )
{
    stage_list[i]->do()
}
```

Our solution applies the following pattern:

```
template<class typelist>
struct ff_pipeline
{
    std::vector<void*> stages;

    void do()
    {
        do_aux<
            0,
            boost::mpl::size<typelist>::value
        >::do(stages);
    }

    template<int I, int N>
    struct do_aux
    {
        static void do(std::vector<void*>& stages)
        {
            static_cast<
                typename boost::mpl::at<
                    typelist,
                    boost::mpl::int_<I>
                >::type*>*>(&stages)[I]->do();
        }
    };
};
```

```

        >::type
        >(stages[I])->do();

        do_stages_aux<
            I+1,
            boost::mpl::size<typelist>::value
        >::do(stages);
    }
};

template<int N>
struct do_aux<N, N> /* specialization */
{
    static void do(std::vector<void*>&) { }
};

/* ... */
};

```

Our solution based on recursive template instantiation [1]. The struct `do_aux` has two template arguments: `I` is the loop variable, and `N` is the size of `typelist`. This struct has a static member function called `do` which casts the corresponding stage to the proper type and invokes its `do` member function. After that, it instantiates itself with `I+1` and `N`, and calls its `do` member function recursively. When `I` equals to `N`, the specialized version of `do_aux` is selected. While its static member function does nothing, the recursion will be stopped.

The computation starts by the member function `do` of struct `ff_pipeline`. It instantiates the struct `do_aux` with 0 and the size of the `typelist` and invokes its static member function. The entities in scope `boost::mpl` are described in [20].

The example below shows how to create pipeline in FastFlow:

```

struct Stage1 : ff_node<Stage1, my_data_type>;
struct Stage2 : ff_node<Stage2, my_data_type>;
//...
typedef boost::mpl::vector<Stage1, Stage2>::type stage_types;
ff_pipeline<stage_types, my_data_type>* pipe =
    new ff_pipeline<stage_types, my_data_type>();
pipe->add_stage(new Stage1());
pipe->add_stage(new Stage2());

```

The structs `Stage1` and `Stage2` derived from `ff_node` are the stages of the pipeline and the `stage_types` refers to their type. The `ff_pipeline` is instantiated with that. In the last two rows, that stages are added.

While a template argument of `ff_pipeline` refers to the type of the stages, their structure must be known at compilation time. Therefore it introduces a restriction of the usage. However, in most cases the structure of the pipeline is defined at compilation time, thus it is not a real limitation.

## 6 Buffers

The original version of FastFlow can use either circular or non-circular buffer as a channel between worker threads. These buffers store `void*` to be able to handle different types of data. By default the non-circular buffer is used, however, the programmer can select the other one by defining preprocessor macro `FF_BOUNDED_BUFFER`. The following preprocessor code snippet shows how the proper buffer is chosen.

```

#if defined(FF_BOUNDED_BUFFER)
#define FFBUFFER SWSR_Ptr_Buffer

```



```
#else // non-circular buffer
#define FFBUFFER uWSR_Ptr_Buffer
#endif
```

Later the `FFBUFFER` macro is used as a buffer type.

The idea of writing preprocessor macros is against the modern development approaches in C++. The code will be less readable, harder to understand, and the preprocessor does not know the syntax of the language. The possible errors are figured out only at compile time (after preprocessing) and the line numbers of the errors (presented by the compiler) may be confusing, because they refer to the place where the preprocessor macros are applied, and not where they are defined.

On the other hand, with our type-safe solution the buffers need to store a pointer to a proper type and not to `void*`, thus they need to be template. Defining template buffers with preprocessor macros makes the code more messy.

In our solution we avoid using preprocessor macros to defining the buffers. Instead we add an extra `boolean` template argument to `ff_farm` called `is_circular`, to indicate which kind of buffer to be to use. The `true` value means we want to use circular buffer, and the `false` means we chose the non-circular one. The default value of that template argument is `false`.

At the definition of buffer we apply that template argument as a condition on a template metaprogram `if_` [1], which chooses the proper buffer type during compilation. The following code snippet shows this, where `T` is the type being stored in the buffer:

```
if_<is_circular,
    SWSR_Ptr_Buffer<T>,
    uWSR_Ptr_Buffer<T> >::type buffer;
```

The definition of `if_` is the following:

```
template< bool condition, class then_, class else_ >
struct if_
{
    typedef then_ type;
};

template< class then_, class else_ >
struct if_<false>
{
    typedef else_ type;
}
```

If the argument `condition` is true, the general `if_` structure is chosen to be instantiate, thus the `type` refers to template type argument `then_`. If the argument `condition` is false the partially specialized `if_` structure is chosen in the instantiation process, thus `type` defines the template type argument `else_`.

## 7 Example

In this chapter we show the example of chapter 2 rewritten with templates. Each class derived from `ff_node` instantiated with itself and the task type `int`. The function `svc` in each classes has a properly typed pointer, instead of `void*`. The `main` function creates the `farm`, the `emitter`, `worker` and `collector` classes, and starts computing.

```

// the generic worker
struct Worker: ff::ff_node<Worker, int>
{
    int * svc( int * task )
    {
        std::cout << "Worker "
                    << ff_node<Worker, int>::get_my_id()
                    << " received task " << *task << "\n";
        return task;
    }
};

// the gatherer filter
struct Collector: ff::ff_node<Collector, int>
{
    int * svc( int * task )
    {
        if ( *task == -1 ) return NULL;
        return task;
    }
};

// the load-balancer filter
struct Emitter: ff::ff_node<Emitter, int>
{
    Emitter( int max_task ):ntask( max_task ) {}

    int * svc( int * )
    {
        int * task = new int( ntask );
        --ntask;
        if ( ntask<0 ) return NULL;
        return task;
    }
private:
    int ntask;
};

int main()
{
    ff_farm<Emitter, Worker, Collector, int> farm;
    Emitter e( TASK_NUM );
    Collector c;
    std::vector<ff::ff_node<Worker, int> *> workers;
    for( int i = 0; i < WORKER_NUM; ++i )
        workers.push_back( new Worker() );
    farm.add_emitter( &e );
    farm.add_workers( workers );
    farm.add_collector( &c );

    farm.run_and_wait_end();
}

```

## 8 Conclusion

Multicore programming is an interesting new way of programming. Many C++ libraries support multicore programming, however, Fastflow is one of the most promising ones.

In this paper we pointed out a weakness of the FastFlow multicore programming library of the C++ programming language. The original version of FastFlow uses `void*` to communicate the between threads. While this method provides a simple solution to pass arbitrary data, it is not typesafe and increases the chance of make mistakes.

We presented our solution to overcome of this problem. We introduced template arguments to deal with the type of arguments and static polymorphism is applied insted of dynamic polymorphism.

In this paper we do not deal with the speedup of FastFlow because it would

be misleading. We have eliminated virtuality of a very important method in the library, but we have not changed anything in the multicore functionality. The typical conception of efficacy of a multicore library is based on the features of multicore programming and it is unchanged in this case.

Our approach also shows why generative technologies are strengthened among multicore programming, too.

## Acknowledgement

The Project is supported by the European Union and co-financed by the European Social Fund (grant agreement no. TÁMOP 4.2.1./B-09/1/KMR-2010-0003).

## References

- [1] Abrahams, D., Gurtovoy, A.: “C++ Template Metaprogramming: Concepts, Tools, and Techniques from Boost and Beyond”, Addison-Wesley (2004).
- [2] Aldinucci, M., Danelutto, M., Meneghin, M., Kilpatrick, P., Torquati, M.: *Efficient streaming applications on multi-core with FastFlow: the biosequence alignment test-bed*, in Proc. of Intl. Parallel Computing (PARCO) 2009.
- [3] Aldinucci, M., Meneghin, M., Torquati, M.: *Efficient smith-waterman on multi-core with FastFlow*, In Proc. of Intl. Euromicro PDP 2010: Parallel Distributed and network-based Processing, IEEE, pp. 195–199.
- [4] Aldinucci, M., Ruggieri, S., Torquati, M.: *Porting Decision Tree Algorithms to Multicore using FastFlow*, in Proc. of European Conference in Machine Learning and Knowledge Discovery in Databases (ECML PKDD), LNCS **6321**, pp. 7–23.
- [5] Alexandrescu, A.: “Modern C++ Design”, Addison-Wesley (2001).
- [6] Bishof, H., Gorlatsch, S.: *Generic parallel programming using C++ templates and skeletons*, in Proc. of Domain-Specific Program Generation (International Seminar), 2003, Revised Papers, LNCS **3016**, pp. 107–126.
- [7] Czarnecki K., Eisenecker, U. W.: “Generative Programming: Methods, Tools and Applications”, Addison-Wesley (2000).
- [8] Dagum, L., Menon, R.: *OpenMP: An industry-standard API for shared-memory programming*, “IEEE Computational Science and Engineering” **5** (1998), pp. 46–55.
- [9] Duret-Lutz, A., Géraud, T., Demaille, A.: *Design patterns for generic programming in C++*, In Proc. of the 6th USENIX Conference on Object-Oriented Technologies and Systems (COOTS), pp. 189–202.
- [10] Harrison, N., Meiners, J. H.: *The dynamics of changing dynamic memory allocation in a large-scale C++ application*, In Companion To the 21st ACM SIGPLAN Conference on Object-Oriented Programming Systems, Languages, and Applications (Portland, Oregon, USA, October 22 - 26, 2006) (OOPSLA, 2006), pp. 866–873.
- [11] Leiserson, C. E.: *The Cilk++ Concurrency Platform*, in Proc. of 46th Annual Design Automation Conference 2009, pp. 26–31.
- [12] Pataki, N.: *C++ Standard Template Library by safe functors*, in Proc. of the 8th Joint Conference on Mathematics and Computer Science (MaCS 2010), Selected Papers, pp. 363–374.
- [13] Pheatt, C.: *Intel Threading Building Blocks*, “Journal of Computing Sciences in Colleges” **23(4)**, pp. 298–298.
- [14] Reinders, J.: “Intel Threading Building Blocks”, O’Reilly (2007).
- [15] Stroustrup, B.: “The C++ Programming Language”, Special Edition, Addison-Wesley (2000).
- [16] Szűgyi, Z., Pataki, N.: *Sophisticated Methods in C++*, In Proc. of International Scientific Conference on Computer Science and Engineering (CSE 2010), pp. 93–100.

- [17] Szűgyi, Z., Pataki, N.: *A More Efficient and Type-Safe Version of FastFlow*, in Proc. of Workshop on Generative Technologies (WGT 2011), pp. 24–37.
- [18] Szűgyi, Z., Sinkovics, Á., Pataki, N., Porkoláb, Z.: *C++ Metastring Library and its Applications*, In Proc. of Generative and Transformational Techniques in Software Engineering 2009, LNCS **6491**, pp. 461–480.
- [19] Szűgyi, Z., Török, M., Pataki, N.: *Towards a Multicore C++ Standard Template Library*, in Proc. of Workshop on Generative Technologies (WGT 2011), pp. 38–48.
- [20] [http://www.boost.org/doc/libs/1\\_45\\_0/libs/mpl/doc/index.html](http://www.boost.org/doc/libs/1_45_0/libs/mpl/doc/index.html)