



ELSEVIER

Available online at [www.sciencedirect.com](http://www.sciencedirect.com)

ScienceDirect

---

**Electronic Notes in  
Theoretical Computer  
Science**

---

Electronic Notes in Theoretical Computer Science 238 (2009) 71–82

[www.elsevier.com/locate/entcs](http://www.elsevier.com/locate/entcs)

# Towards the Modularization of C++ Concept Maps

Zalán Szűgyi<sup>1</sup> Ádám Sipos<sup>2</sup> Zoltán Porkoláb<sup>3</sup>*Department of Programming Languages and Compilers  
Eötvös Loránd University  
Budapest, Hungary*

---

## Abstract

*Concept* is a new and powerful language element being introduced in the next C++ standard. With the help of concepts we can define the properties a template requires from its type arguments. If a type does not fulfill the requirements syntactically, but semantically, the connection must be declared with the use of a `concept_map`. Often the description of the semantic matching results in long codes that need to be modularized. In this paper we present an extension to the `concept_map` language constructs that enables this transformation. We introduce the well-known `public`, `protected`, `private` class arrangement scheme into `concept_maps`. We present our preprocessor, that transforms the modularized code into regular code to be compiled by ConceptGCC, the experimental Concept C++ compiler.

*Keywords:* C++, ConceptC++

---

## 1 Introduction

*Generative programming*[2] is a new programming paradigm, focusing on program transformation, code generation, and the development of highly reusable libraries. The C++ programming language supports generative programming by the use of its generic construct, the *template*[11]. Templates in C++ are *unconstrained*, so it is possible to instantiate a class template or function template with any type argument. The instantiation process is carried out irrespective of whether or not the argument satisfies the requirements implicitly set by the template. Thus the relevant compiler errors are triggered too late in the instantiation process causing complicated and difficult to understand error messages. Writing a good documentation and giving informative names to our template's parameters is not a sufficient solution for informing the user of the template's requirements.

---

<sup>1</sup> Email: [lupin@inf.elte.hu](mailto:lupin@inf.elte.hu)

<sup>2</sup> Email: [shp@inf.elte.hu](mailto:shp@inf.elte.hu)

<sup>3</sup> Email: [gds@inf.elte.hu](mailto:gds@inf.elte.hu)

The research area *concept checking* identified this problem and provided library-level solutions [7,8]. A cleverly designed C++ code is capable of deciding at compile-time if two classes are in a parent-child relationship, or if a type is a pointer type, and other properties. The most notable concept checking library containing such constructs is `boost::concept_check` [12]. The `boost::concept_check` library included and unified many previous approaches, and attracted attention to this area. However, the library still lacks a number of functionalities, and proved that a language-based solution is required rather than a library-based one.

A new C++ language construct, the *concept* [3] solves the problem of compile-time type checking at the language level by adding constrained templates to the language, and still keeping all of the advantages of the C++ template system. A `concept` defines a set of requirements on one or more types. These requirements describe the *interface* a type needs to offer to meet the template's precondition. The interface is described with *function signatures* and *associated types*. If a type satisfies all the requirements of a concept we say that type *models* this concept.

In many cases a class does not syntactically match the requirements of a `concept`. For example the `concept` may require that the *plus* (+) operator be defined on the argument type [13]. In the case of a `Color` type, however, the *plus* operation may be modelled with the method `mix()` mixing two colors. This *semantic* matching can be described with a `concept_map`.

Originally `concept_maps` were intended to be short code parts with function signatures and short function bodies. However, nowadays `concept_maps` are being used for a number of new and complex purposes, like serving as *adaptors* between a class and an interface (see the `std::vector – Stack` example in Section 3). Thus the description of a semantic match may result in a longer and more complicated program code. Currently this whole code must be written either in the function body or with the help of other function calls. Avoiding repetition of common code fragments in `concept_maps` also requires the implementation of helper functions. However, currently such functions have to be declared and implemented outside of the `concept_map`. Thus, helper functions regularly declared as free functions or placed in ad-hoc implementational namespaces, leading to scattering helper functions in the source and seriously decrease the quality of the code.

This phenomenon is in sharp contrast with other constructs of the C++ language. Classes, for instance, support better modularization: private and public visibility distinguish helpers from interface functions, in the same time avoid code scattering grouping them into one code unit.

In this paper we present an approach for modularizing `concept_maps` thus creating code that is easier to understand, and maintain. We present our preprocessor implementing the approach by transforming the modularized code into C++ code accepted by the experimental ConceptGCC compiler.

The rest of the paper is organized as follows. In Section 2 we discuss the `concept` and `concept_map` language constructs. In Section 3 we describe the `concept_map` modularization problem to be addressed in the paper. Section 4 discusses our solution for the problem, and the preprocessor implementing our approach. In

Section 6 we conclude our paper and discuss future work.

## 2 Concepts

The language construct `concept` is essentially a listing of function signatures and associated types that we intend to use in our template. For example, the following concept requires that the equality `operator==` is defined for types `T` and `U`.

```
concept EqualityComparable<typename T, typename U = T>
{
    bool operator==(T a, U b);
};
```

Every type whose equality operator is defined models this concept. The equality operator can be defined as a member or a non member function. We can refine an existing concept by adding new requirements for it. For example we add a requirement for the less than (`<`) operator to the previous concept.

```
concept Comparable<typename T, typename U = T>
    : EqualityComparable<T,U>
{
    bool operator<(T a, U b);
};
```

There is another way to create a concept, using the keyword `auto` before keyword `concept`. Consider the following example:

```
auto concept EqualityComparable<typename T, typename U = T>
{
    bool operator==(T a, U b);
};
```

In the first two cases it must be explicitly declared how the type satisfies the requirements, and in the third case this is not necessary. The `concept_map` is the language construct used for describing the semantical connection between the type and the concept. A type can satisfy the requirements *syntactically* if it contains the functions and types that the concept requires. The type satisfies the requirements *semantically* if even though the type does not have all of these functions, the usage of these functions is semantically rational. Let us consider the following example.

```
struct Person
{
    int id;
    string name;
};
```

There is no equality operator for `Person`, but we can consider two persons equal when their `id`-s are equal. Using `concept_map` it is possible to explicitly declare this semantic equality.

```
concept_map EqualityComparable<Person>
{
    bool operator==(Person a, Person b)
    {
        return a.id == b.id;
    }
};
```

If the type syntactically satisfies the requirements (for example `int` has an equality operator) we can just write the following:

```
concept_map EqualityComparable<int>
{
}
```

Note that in case the `concept` is declared as an `auto concept` we do not need to make this mapping.

In the following example we present how to write a constrained function template. The following function utilizes the equality operator of type `T` and declares this with the `requires` keyword.

```
template<typename T>
requires EqualityComparable<T>
void f(const T& t1, const T& t2)
{
    ... if(t1 == t2) ...
}
```

Function `f` can be instantiated only with types satisfying the `EqualityComparable` concept.

### 3 Transforming the `concept_map`

A `concept_map` declares explicitly how a type satisfies the requirements of a concept. In the context of one concept, with `concept_map` we can redefine some member functions of a type, or add some new ones without doing any modification inside the type. For example, in the previous chapter the `Person struct` has no equality operator, but with a `concept_map` we can define one. Thus `Person` can be passed to any class template or function template requiring an `EqualityComparable` argument. The `concept_map` makes it easier to solve certain problems. Let us suppose we need an adaptor for the container `std::vector`. An obvious though tedious approach is to write a wrapper class `Stack` for the container. However, we can utilize a `concept_map` and essentially rename the container member functions. In the following example, we demonstrate how to create a stack using `std::vector`[\[4\]](#).

```
concept Stack<typename X>
{
    typename value_type;
```

```

void push(X&, const value type&);
void pop(X&);
value type top(const X&);
};

```

```

template<typename T>
concept_map Stack<std::vector<T> >
{
    typedef T value type;
    void push(std::vector<T>& v, const T& x) { v.push back(x); }
    void pop(std::vector<T>& v) { v. pop back(); }
    T top(const std::vector<T>& v) { return v. back(); }
};

```

With the `concept_map` we have described how the `std::vector` can be regarded as a `Stack`. Now we can instantiate every class template or function template with `std::vector`, even if the template argument is required to fulfill the `Stack` concept. In these examples the code of the `concept_maps` is only two or three lines long. However, many times with complex concept mappings the code describing the semantical matching may be more complicated and longer. In our next example we present how a to simulate an aspect-oriented feature [6] with the help of `concept_map`. Let us suppose we want to log data before and after calling the `f()` and `g()` functions of type `MyType`. Our logging procedure consists of opening a network connection, carrying out file I/O operations, and ending the connection.

```

concept LogConcept<typename C>
{
    void f(C&);
    void g(C&);
}

concept_map LogConcept<MyType>
{
    void f(MyType& c)
    {
        // open network connection
        // open logfile
        // log into file
        // close logfile
        // close network connection
        // error handling
        c.f();
        // open network connection
        // open logfile
        // log into file
    }
}

```

```

    // close logfile
    // close network connection
    // error handling
}

void g(MyType& c)
{
    // open network connection
    // open logfile
    // log into file
    // close logfile
    // close network connection
    // error handling
    c.g();
    // open network connection
    // open logfile
    // log into file
    // close logfile
    // close network connection
    // error handling
}
};

```

The logging part of this `concept_map` contains several redundant lines of code. On one hand this leads to long function bodies in the `concept_map`. On the other hand the actual original function calls (`c.f()` and `c.g()`) will be lost in the long code. It may be useful to trigger a separate helper function call at these points, however, there is no well situated place to put these functions. It is possible to put the helper functions into the global namespace, but if they are related only to the `concept_map` this solution will be confusing and makes our code harder to understand.

In the following we describe our remedy for these problems.

## 4 Modularization

Our approach is to introduce helper functions into `concept_maps`. These functions do not express some requirement for the template argument, but rather are called from bodies of the various functions of the `concept_map`, just as if they were ordinary C++ functions. The `concept_map` structure is extended with `private` and `protected` and `public` parts as done in `class` or `struct` constructs. Functions defined as `private` or `protected` are hidden outside the `concept_map`. The `public` keyword is optional, as all requirements (function signatures) are public by default. Consequently, our approach is backward compatible with the current `concept` mechanism: if none of the previous visibility keywords are used, the `concept_map` has its regular form, and can be compiled. At this stage of the research the `protected` and `private` keywords provide the same functionality. One of our most important

future plans is to introduce some refined meaning for **protected**.

Let us consider the following example. We modified the example about logging in the previous chapter using the **public** and **private** keywords.

```
concept LogConcept<typename C>
{
    void f(C&);
    void g(C&);
}

concept_map LogConcept<MyType>
{
    public:
        void f(MyType& c)
        {
            log();
            c.f();
            log();
        }

        void g(MyType& c)
        {
            log();
            c.g();
            log();
        }

    private:
        void log()
        {
            // open network connection
            // open logfile
            // log into file
            // close logfile
            // close network connection
            // error handling
        }
};
```

In the function bodies of **f()** and **g()** we have substituted the redundant logging procedure with simple function calls. Consequently our **concept\_map** is shorter, more readable, avoids code duplication, and is easier to maintain. The **log()** function declared as **private** will be unreachable from the outside of the **concept\_map**, thus only member functions will be able to use it. This extended notation of **concept\_maps** is not a valid ConceptC++ construct, thus a compiler extension is required to be able to handle the notation.

The only compiler currently implementing **concepts** for C++ is *Concept-GCC*[13]. We are not intending to create a new **concept** compiler, rather to write a preprocessor transforming our code into valid ConceptC++ code and still keeping our new structural advantages. The idea is to put the **private** and **protected** parts of the **concept\_map** to a unique **namespace** and put the namespace scope before the function calls in the **public** part. We have implemented a precompiler-like program to carry out this transformation. After the transformation we get the following code:

```
namespace unique_ns_name
{
    void log()
    {
        ...
    }
}

concept_map LogConcept<MyType>
{
    void f(MyType& c)
    {
        unique_ns_name::log();
        c.f();
        unique_ns_name::log();
    }
    void g(MyType& c)
    {
        unique_ns_name::log();
        c.g();
        unique_ns_name::log();
    }
};
```

As seen in example 9 the **concept\_map** itself can be template. In this case the template argument may appear inside the helper functions. Transforming this kind of **concept\_maps** is not so simple as it was in the previous example. When the helper function definition is moved from the body of the **concept\_map** to the unique namespace, the template arguments of the **concept\_map** became unknown types. Thus it is necessary to transform these helper functions to template functions with the same template arguments as the **concept\_map**. It is necessary to instantiate these helper functions explicitly in the **public** part of the **concept\_map**, because the template arguments may not be part of the function parameter list. In the following example we create a new **concept\_map** (which contains helper functions) to concept **Stack** defined in example 9.

```
template<typename T>
```



```

concept_map Stack<std::vector<T> >
{
    public:
        typedef T value type;
        void push(std::vector<T>& v, const T& x)
        {
            f();
            v.push_back(x);
            g();
        }
        void pop(std::vector<T>& v) { v.pop_back(); }
        T top(const std::vector<T>& v) { return v.back(); }
    private:
        void f() { ... T t; ... do something before push ... }
        void g() { ... T t; ... do something after push ... }
};

```

After the transformation we get the following code:

```

namespace unique_ns_name
{
    template<typename T>
    void f() { ... T t; ... do something before push ... }

    template<typename T>
    void g() { ... T t; ... do something after push ... }
}

template<typename T>
concept_map Stack<std::vector<T> >
{
    typedef T value type;
    void push(std::vector<T>& v, const T& x)
    {
        unique_ns_name::f<T>();
        v.push_back(x);
        unique_ns_name::g<T>();
    }
    void pop(std::vector<T>& v) { v.pop_back(); }
    T top(const std::vector<T>& v) { return v.back(); }
};

```

In some cases the member functions are so long that it is not practical to store them within the body of the `concept_map` (similarly to the case of C++ `structs` and `classes` and their member functions). Our approach provides a solution for these cases, as it is possible to define the functions outside the type definition. When

the `concept_map` is not a template the helper function can be defined in a separate file. It is necessary to put the outside definition into the same namespace as the `concept_map`. In the following example we implement the logger functions outside the `concept_map`.

```
concept_map LogConcept<MyType>
{
    public:
        void f(MyType& c)
        {
            log();
            c.f();
            log();
        }
        void f(MyType& c)
        {
            log();
            c.g();
            log();
        }
    private:
        void log();
}

// function can be placed into a separate file
void concept_map LogConcept<MyType>::log()
{
    ...
}
```

After the transformation we get the following code:

```
// Forward declaration is needed because
// the definition of the function can be in a separate file
namespace unique_ns_name
{
    void log();
}

concept_map LogConcept<MyType>
{
    void f(MyType& c)
    {
        unique_ns_name::log();
        c.f();
        unique_ns_name::log();
    }
}
```

```

    }
}

// can be placed into a separate file
namespace unique_ns_name
{
    void log()
    {
        ...
    }
}

```

## 5 Related work

A strong restriction on `concept_maps` is that it must be in the same namespace as the original `concept`. If we intend to use a third-party library (as in [5]), we must put our own `concept_maps` into the library's namespace. This is not a clean approach, and may obviously cause problems. In [9] the authors present a detailed description of the problem and also a solution. Our approach can fully cooperate with their proposed solution, as these two problems are independent of each other.

## 6 Conclusion

In this paper we have presented a new approach for modularizing the new C++ language construct `concept_map`. The keywords `private`, `public`, and `protected` are now available for use in `concept_maps`, resulting in more maintainable code. Our preprocessor is capable of recognizing `concept_maps`, and transforming the contained code into simple ConceptGCC code. Our solution is backward compatible with the proposed C++ Standard, i.e. omitting keywords we get semantically equivalent code with the Concepts proposal.

## References

- [1] Bourdev, L. and H. Jin: *Generic Image Library*, <http://adobe.com/gil/> 2006. opensource.
- [2] Czarnecki, K. and U. W. Eisenecker: *Generative Programming: Methods, Tools and Applications*, Addison-Wesley 2000.
- [3] Douglas, G., J. Järvi, J. G. Siek, G. D. Reis, B. Stroustrup and A. Lumsdaine: *Concepts: Linguistic Support for Generic Programming in C++*, "Proceedings of the 2006 ACM SIGPLAN conference on Object-oriented programming, systems, languages, and applications (OOPSLA'06)" (2006)
- [4] Douglas, G. and B. Stroustrup: *Concepts (Revision 1)*, 2006.
- [5] Järvi, J., M. A. Marcus and J. N. Smith: *Library Composition and Adaptation using C++ Concepts*, In Proc. of the 6th international Conference on Generative Programming and Component Engineering, pp. 73-82, 2007.
- [6] Kiczales, G., J. Lamping, A. Menhdhekar, C. Maeda, C. Lopes, J-M. Loingtier, and J. Irwin: *Aspect-Oriented Programming*, Aksit et al: "Proceedings European Conference on Object-Oriented Programming 1241" (1997), Springer-Verlag, pp. 220-242, 1997.

- [7] McNamara, B., Smaragdakis, Y.: *Static Interfaces in C++*, In First Workshop on C++ Template Metaprogramming, October 2000, Erfurt.
- [8] Siek, J., Lumsdaine, A.: *Concept Checking: Binding Parametric Polymorphism in C++*, In First Workshop on C++ Template Metaprogramming, October 2000, Erfurt.
- [9] Siek, J.: *Scoped Concept Maps*, C++ Standards Committee Paper.
- [10] Siek, J., L.-Q. Lee and A. Lumsdaine: *The Boost Graph Library: User Guide and Reference Manual*, Addison-Wesley Longman Publishing Co., Inc., Boston, MA, USA, 2002.
- [11] Stroustrup, B.: *The C++ Programming Language Special Edition*, Addison-Wesley 2000.
- [12] *The Boost Concept Checking Library*. [http://www.boost.org/libs/concept\\_check/concept\\_check.htm](http://www.boost.org/libs/concept_check/concept_check.htm)
- [13] *ConceptGCC BoostCon Edition*. <http://www.generic-programming.org/software/ConceptGCC/>