

Full Abstraction Without Synchronization Primitives

Andrzej S. Murawski¹

*Oxford University Computing Laboratory
Wolfson Building, Parks Road, Oxford OX1 3QD, UK*

Abstract

Using game semantics, we prove a full abstraction result (with respect to the may-testing preorder) for Idealized Algol augmented with parallel composition (IA^{\parallel}). Although it is common knowledge that semaphores can be implemented using shared memory, we find that semaphores do not extend IA^{\parallel} conservatively. We explain the reasons for the mismatch.

Keywords: Shared-Variable Concurrency, Mutual Exclusion, Full Abstraction, Game Semantics

1 Introduction

The mutual exclusion problem asks one to find sections of code that will allow two threads to share a single-use resource without conflict. It turns out that shared memory (with atomic reads and writes) can be used to solve it without any additional synchronization primitives. A typical solution consists of two sections of code (called *entry* and *exit* protocols respectively) that each of the two processes can use to enter and exit their designated critical sections respectively.

Quite a collection of trial solutions have been shown to be incorrect and at some moment people that had played with the problem started to doubt whether it could be solved at all.

So writes Dijkstra [3] about early attempts to attack the problem. He credits Dekker with the first correct solution, which was later simplified by several other authors.

¹ Supported by an EPSRC Advanced Research Fellowship (EP/C539753/1).

Peterson’s tie-breaker algorithm [10], reproduced below, was particularly elegant.

<pre>/* Entry Code 1 */ Q[1] := 1; turn := 1; while (Q[2] and (turn = 1)) do skip; /* Exit Code 1 */ Q[1] := 0;</pre>	<pre>/* Entry Code 2 */ Q[2] := 1; turn := 2; while (Q[1] and (turn = 2)) do skip; /* Exit Code 2 */ Q[2] := 0;</pre>
--	--

Solutions to the two-process case were subsequently generalized to n processes (Lamport’s bakery algorithm [7] is one of the simplest). Although the results demonstrated that, from a theoretical point of view, the sharing of memory was sufficient to enforce mutual exclusion, they were considered unsatisfactory from the conceptual and implementation-oriented points of view. The intricacy of interactions generated by the code was judged to obscure the purpose it was supposed to serve and the “busy-waiting” involved looked wasteful. This motivated the introduction of semaphores [3], a synchronization construct on a higher level than memory reads and writes.

In this paper we would like to focus on the expressive power of semaphores in the setting of shared-variable higher-order concurrency and contextual testing. We consider a variant IA of Reynolds’ Idealized Algol [11] augmented with parallel composition, referred to as IA^{\parallel} , and prove an inequational full abstraction result for the induced notion of contextual *may-testing*. The result is obtained using game semantics by uncovering a preorder on strategies, founded on a notion reminiscent of racing computations.

Contrary to what the various mutual-exclusion algorithms might suggest, we find that there are strategies corresponding to programs with semaphores, which do not correspond to any IA^{\parallel} -terms. What is more, we can identify a game-semantic closure property enjoyed by all strategies corresponding to IA^{\parallel} -terms, which may fail in the presence of semaphores. This makes it possible to apply our model to the semantic detection of the “need for semaphores”. As for contextual may-approximation and may-equivalence, we show that IA^{\parallel} extended with semaphores does *not* constitute a conservative extension of IA^{\parallel} . We conclude by relating the apparent mismatch to non-uniformity of mutual-exclusion algorithms based on shared memory alone.

From the game-semantic perspective, our results demonstrate that a language without semaphores is considerably more difficult to handle than one incorporating them. So, the addition of communication primitives to a language can lead to cleaner mathematical structure.

2 Idealized Algols

We shall be concerned with parallel extensions of Reynolds’ Idealized Algol [11], which has become the canonical blueprint for synthesizing imperative and functional

TYPES

$$\beta ::= \mathbf{com} \mid \mathbf{exp} \mid \mathbf{var} \qquad \theta ::= \beta \mid \theta \rightarrow \theta$$

TERMS

$$\begin{array}{c} \overline{\Gamma \vdash \mathbf{skip} : \mathbf{com}} \qquad \overline{i \in \mathbb{N}} \quad \Gamma \vdash i : \mathbf{exp} \qquad \overline{\Gamma, x : \theta \vdash x : \theta} \\[10pt] \frac{\Gamma \vdash M_1 : \mathbf{com} \quad \Gamma \vdash M_2 : \beta}{\Gamma \vdash M_1; M_2 : \beta} \qquad \frac{\Gamma \vdash M_1 : \mathbf{exp} \quad \Gamma \vdash M_2 : \mathbf{exp}}{\Gamma \vdash M_1 \oplus M_2 : \mathbf{exp}} \\[10pt] \frac{\Gamma \vdash M : \mathbf{exp} \quad \Gamma \vdash N_0 : \theta \quad \Gamma \vdash N_1 : \theta}{\Gamma \vdash \mathbf{if } M \mathbf{ then } N_1 \mathbf{ else } N_0 : \theta} \\[10pt] \frac{\Gamma \vdash M : \mathbf{var}}{\Gamma \vdash !M : \mathbf{exp}} \quad \frac{\Gamma \vdash M : \mathbf{var} \quad \Gamma \vdash N : \mathbf{exp}}{\Gamma \vdash M := N : \mathbf{com}} \quad \frac{\Gamma, x : \mathbf{var} \vdash M : \beta}{\Gamma \vdash \mathbf{newvar } x \mathbf{ in } M : \beta} \\[10pt] \frac{\Gamma \vdash M : \mathbf{exp} \rightarrow \mathbf{com} \quad \Gamma \vdash N : \mathbf{exp}}{\Gamma \vdash \mathbf{mkvar}(M, N) : \mathbf{var}} \\[10pt] \frac{\Gamma \vdash M : \theta \rightarrow \theta' \quad \Gamma \vdash N : \theta}{\Gamma \vdash MN : \theta'} \quad \frac{\Gamma, x : \theta \vdash M : \theta'}{\Gamma \vdash \lambda x^\theta. M : \theta \rightarrow \theta'} \quad \frac{\Gamma, x : \theta \vdash M : \theta}{\Gamma \vdash \mu x^\theta. M : \theta} \end{array}$$

Fig. 1. Syntax of IA

programming. The particular variant of Idealized Algol, presented in Figure 1 and henceforth referred to as **IA**, is known in the literature as Idealized Algol with active expressions [1]. This paper is primarily devoted to **IA** extended with the parallel composition operator \parallel . It enters the syntax through the following typing rule.

$$\frac{\Gamma \vdash M_1 : \mathbf{com} \quad \Gamma \vdash M_2 : \mathbf{com}}{\Gamma \vdash M_1 \parallel M_2 : \mathbf{com}}$$

We shall write \mathbf{IA}^\parallel to denote the extended language. Our main goal will be to arrive at a fully abstract model for contextual approximation and equivalence induced by \mathbf{IA}^\parallel . In particular, we would like to understand how the addition of semaphores to \mathbf{IA}^\parallel affects the two. To that end, we consider yet another prototypical language, called **PA**, which is \mathbf{IA}^\parallel extended with semaphores. We give the syntax of **PA** in Figure 2. We assume that semaphores and variables are initialized to “available” and 0 respectively.

Remark 2.1 *PA was introduced in [4]. It is closely related to Brookes’ Parallel Algol [2], which, in contrast to PA, represents the coarse-grained approach to enforcing atomicity. Parallel Algol contains the **await** M **then** N construct which executes the guard M as an atomic action and, if the guard is true, N is run immediately afterwards, also as an indivisible operation. PA and Parallel Algol appear equi-expressive. Clearly, semaphores can be implemented using **await** - **then** - and ordinary variables. A translation in the other direction is also possible, for exam-*

TYPES

$$\beta ::= \mathbf{com} \mid \mathbf{exp} \mid \mathbf{var} \mid \mathbf{sem} \qquad \theta ::= \beta \mid \theta \rightarrow \theta$$

TERMS

All rules defining \mathbf{IA}^{\parallel} plus the following ones

$$\frac{\Gamma \vdash M : \mathbf{sem}}{\Gamma \vdash \mathbf{grab}(M) : \mathbf{com}} \qquad \frac{\Gamma \vdash M : \mathbf{sem}}{\Gamma \vdash \mathbf{release}(M) : \mathbf{com}}$$

$$\frac{\Gamma, x : \mathbf{sem} \vdash M : \beta}{\Gamma \vdash \mathbf{newsem} \, x \, \mathbf{in} \, M : \beta} \qquad \frac{\Gamma \vdash M : \mathbf{com} \quad \Gamma \vdash N : \mathbf{com}}{\Gamma \vdash \mathbf{mksem}(M, N) : \mathbf{sem}}$$

Fig. 2. Syntax of PA

ple, in the style of the encoding of Parallel Algol into the π -calculus [12]. We use PA because the game semantics we rely on is better suited to modelling fine-grained concurrency and **await** would have had to be interpreted indirectly by translation.

For a closed \mathbf{IA} -, \mathbf{IA}^{\parallel} -, PA-term $\vdash M : \mathbf{com}$ we shall write $M \Downarrow$ iff there exists a terminating run of M (the reduction rules are routine and can be found, for example, in [4]). Note that our notion of termination is *angelic*. Accordingly, the notions of contextual approximation and equivalence considered here will be consistent with *may-testing* and will not take the possibility of deadlock/divergence into account.

Definition 2.2 Let $\Gamma \vdash M_1, M_2 : \theta$ be \mathbf{IA}^{\parallel} -terms. $\Gamma \vdash M_1 : \theta$ is said to **contextually approximate** $\Gamma \vdash M_2 : \theta$ (written $\Gamma \vdash M_1 \sqsubseteq_{\mathbf{IA}^{\parallel}} M_2 : \theta$), if, and only if, for any \mathbf{IA}^{\parallel} -context $C[-]$ such that $\vdash C[M_i] : \mathbf{com}$ ($i = 1, 2$), $C[M_1] \Downarrow$ implies $C[M_2] \Downarrow$. Further, $\Gamma \vdash M_1 : \theta$ and $\Gamma \vdash M_2 : \theta$ are **contextually equivalent** (written $\Gamma \vdash M_1 \cong_{\mathbf{IA}^{\parallel}} M_2$) if each contextually approximates the other.

Analogously, one can define contextual approximation (resp. equivalence) using terms and contexts of \mathbf{IA} or PA. We shall write $\sqsubseteq_{\mathbf{IA}}$ (respectively $\cong_{\mathbf{IA}}$) or $\sqsubseteq_{\mathbf{PA}}$ (resp. $\cong_{\mathbf{PA}}$ when referring to them). For example, it can be readily seen that $\sqsubseteq_{\mathbf{IA}^{\parallel}}$ is not a conservative extension of \mathbf{IA} .

Example 2.3 The two \mathbf{IA} -terms

$$\lambda f^{\mathbf{exp} \rightarrow \mathbf{com} \rightarrow \mathbf{com}}. \mathbf{newvar} \, x \, \mathbf{in} \, f(!x)(x := !x + 2)$$

$$\lambda f^{\mathbf{exp} \rightarrow \mathbf{com} \rightarrow \mathbf{com}}. \mathbf{newvar} \, x \, \mathbf{in} \, f(!x)(x := !x + 1; x := !x + 1)$$

are \mathbf{IA} -equivalent, but are not PA-equivalent.

The main result of our paper is an explicit characterization of $\sqsubseteq_{\mathbf{IA}^{\parallel}}$ (Theorem 4.5) in terms of a preorder on strategies. It will allow us to demonstrate that PA is *not* a conservative extension of \mathbf{IA}^{\parallel} .

Example 2.4 In view of the results given below, the simplest example illustrating the non-conservativity of PA with respect to \mathbf{IA}^{\parallel} (as far as contextual approximation

is concerned) are the terms x and $x||x$, where x is a free identifier of type **com**. We shall have $x \sqsubseteq_{\mathbf{IA}^{\parallel}} x||x$ and $x \not\sqsubseteq_{\mathbf{PA}} x||x$.

Informally, x approximates $x||x$ in \mathbf{IA}^{\parallel} , because any successful run of $C[x]$ can be closely followed by that of $C[x||x]$ in which each atomic action of the second x takes place right after the corresponding action of the first x (one keeps on racing the other). In contrast, in \mathbf{PA} , x might be instantiated with code that will try to acquire a semaphore, in which case $x||x$ will not terminate (take, for example, $C[] \equiv \mathbf{newsem} \ S \ \mathbf{in} \ ((\lambda x^{\mathbf{com}}.[]) \ \mathbf{grab}(S)))$.

Similar terms demonstrate that contextual equivalence is not preserved either. We have $(x \ \mathbf{or} \ (x||x)) \cong_{\mathbf{IA}^{\parallel}} (x||x)$, but $(x \ \mathbf{or} \ (x||x)) \not\cong_{\mathbf{PA}} (x||x)$, where $M \ \mathbf{or} \ N$ stands for

$$\mathbf{newvar} \ X \ \mathbf{in} \ ((X := 0 \ || \ X := 1); \ \mathbf{if} \ !X \ \mathbf{then} \ M \ \mathbf{else} \ N).$$

3 Game semantics

\mathbf{IA} and \mathbf{PA} have already been studied using game semantics, in [1] and [4] respectively. The full abstraction results presented therein are particularly elegant, as they characterize $\sqsubseteq_{\mathbf{IA}}$ and $\sqsubseteq_{\mathbf{PA}}$ via (complete-)play containment. Next we shall review the game model of \mathbf{PA} (originally presented in [4]), as our full abstraction result for \mathbf{IA}^{\parallel} will be phrased in terms of strategies from that model. More precisely, we are going to exhibit a preorder, different from inclusion, that will turn out to capture contextual approximation in \mathbf{IA}^{\parallel} . The induced equivalence relation, characterizing $\cong_{\mathbf{IA}^{\parallel}}$, is also different from play equivalence.

Game semantics uses arenas to interpret types.

Definition 3.1 *An arena A is a triple $\langle M_A, \lambda_A, \vdash_A \rangle$, where*

- M_A is a set of moves;
- $\lambda_A : M_A \rightarrow \{O, P\} \times \{Q, A\}$ is a function determining whether $m \in M_A$ is an **Opponent** or a **Proponent** move, a **question** or an **answer**; we write $\lambda_A^{OP}, \lambda_A^{QA}$ for the composite of λ_A with respectively the first and second projections;
- \vdash_A is a binary relation on M_A , called **enabling**, such that $m \vdash_A n$ implies $\lambda_A^{QA}(m) = Q$ and $\lambda_A^{OP}(m) \neq \lambda_A^{OP}(n)$. Moreover, if $n \in M_A$ is such that $m \not\vdash_A n$ for any $m \in M_A$ then $\lambda_A(n) = (O, Q)$.

If $m \vdash_A n$ we say that m enables n . We shall write I_A for the set of all moves of A which have no enabler; such moves are called *initial*. Note that an initial move must be an Opponent question.

In arenas used to interpret base types all questions are initial and all P-moves are answers enabled by initial moves as detailed in the table below, where $m \in \mathbb{N}$.

Arena	O-question	P-answers	Arena	O-question	P-answers
[[com]]	<i>run</i>	<i>done</i>	[[exp]]	<i>q</i>	<i>m</i>
[[var]]	<i>read</i> <i>write(m)</i>	<i>m</i> <i>ok</i>	[[sem]]	<i>grab</i> <i>release</i>	<i>ok^g</i> <i>ok^r</i>

Contexts and function types are modelled with the help of additional constructions on arenas:

$$\begin{aligned}
M_{A \times B} &= M_A + M_B & M_{A \Rightarrow B} &= M_A + M_B \\
\lambda_{A \times B} &= [\lambda_A, \lambda_B] & \lambda_{A \Rightarrow B} &= [\langle \lambda_A^{PO}, \lambda_A^{QA} \rangle, \lambda_B] \\
\vdash_{A \times B} &= \vdash_A + \vdash_B & \vdash_{A \Rightarrow B} &= \vdash_A + \vdash_B + \{ (b, a) \mid b \in I_B \text{ and } a \in I_A \}
\end{aligned}$$

The function $\lambda_A^{PO} : M_A \rightarrow \{O, P\}$ is defined by $\lambda_A^{PO}(m) = O$ iff $\lambda_A^{OP}(m) = P$.

Arenas provide all the details necessary to specify the allowable exchanges of moves. Formally, they will be justified sequences satisfying some extra properties. A *justified sequence* in arena A is a finite sequence of moves of A equipped with pointers. The first move is initial and has no pointer, but each subsequent move n must have a unique pointer to an earlier occurrence of a move m such that $m \vdash_A n$. We say that n is (explicitly) **justified** by m or, when n is an answer, that n **answers** m . If a question does not have an answer in a justified sequence, we say that it is *pending* (or *open*) in that sequence. In what follows we use the letters q and a to refer to question- and answer-moves respectively, o and p to stand for O- and P-moves, and m to denote arbitrary moves.

Not all justified sequences will be regarded as valid. In order to constitute a legal play, a justified sequence must satisfy a well-formedness condition, which reflects the “static” style of concurrency in PA: any process starting sub-processes must wait for the children to terminate in order to continue. In game terms: if a question is answered then all questions justified by it must have been answered earlier (exactly once). This is made precise in the following definition.

Definition 3.2 *The set P_A of plays over A consists of justified sequences s over A satisfying the two conditions below.*

FORK *In any prefix $s' = \dots q \overset{\curvearrowright}{\dots} m$ of s , the question q must be pending before m is played.*

WAIT *In any prefix $s' = \dots q \overset{\curvearrowright}{\dots} a$ of s , all questions justified by q must be answered.*

Note that interleavings of justified sequences are not justified sequences; instead we shall call them *shuffled sequences*. For two shuffled sequences s_1 and s_2 , $s_1 \amalg s_2$ denotes the set of all interleavings of s_1 and s_2 . For two sets of shuffled sequences S_1 and S_2 , $S_1 \amalg S_2 = \bigcup_{s_1 \in S_1, s_2 \in S_2} s_1 \amalg s_2$. Given a set X of shuffled sequences, we define $X^0 = X$, $X^{i+1} = X^i \amalg X$. Then X^* , called *iterated shuffle* of X , is defined to be $\bigcup_{i \in \mathbb{N}} X^i$.

We say that a subset σ of P_A is *O-complete* if $s \in \sigma$ and $so \in P_A$ entail $so \in \sigma$.

Definition 3.3 *A strategy σ on A (written $\sigma : A$) is a prefix-closed and O-complete subset of P_A .*

Strategies $\sigma : A \Rightarrow B$ and $\tau : B \Rightarrow C$ are composed in the standard way, by considering all possible interactions of plays from τ with shuffled sequences of σ^{\oplus} in the shared arena B , and then hiding the B moves.

For modelling concurrent programs, one considers a special class of so-called saturated strategies, which contain all possible (sequential) observations of the relevant (parallel) interactions. Consequently, actions of the environment (O-moves) can always be observed earlier (as soon as they have been enabled), actions of the program can always be observed later (but not later than moves that they justify). To formalize this, for any arena A , one defines a preorder \preceq on P_A as the least reflexive and transitive relation satisfying $s_0 s_1 o s_2 \preceq s_0 o s_1 s_2$ and $s_0 p s_1 s_2 \preceq s_0 s_1 p s_2$ for all s_0, s_1, s_2 . In the above-mentioned pairs of plays, moves on the left-hand side of \preceq are meant to have the same justifiers as on the right-hand side. The two saturation conditions, in various formulations, have a long history in the semantics of concurrency [13,5,6].

Definition 3.4 *A strategy σ is saturated iff $s \in \sigma$ and $s \preceq s'$ imply $s' \in \sigma$.*

Arenas and saturated strategies form a Cartesian closed category \mathcal{G}_{sat} , in which $\mathcal{G}_{\text{sat}}(A, B)$ consists of saturated strategies on $A \Rightarrow B$. The identity strategy is defined by “saturating” the alternating plays $s \in P_{A_1 \Rightarrow A_2}$ in which P “copies” O-moves to the other A -component (formally, for any even-length prefix t of s we have $t \upharpoonright A_1 = t \upharpoonright A_2$). We used A_1 and A_2 to distinguish the two copies of A in the arena $A \Rightarrow A$).

PA-terms $x_1 : \theta_1, \dots, x_n : \theta_n \vdash M : \theta$ can be interpreted in \mathcal{G}_{sat} as strategies in the arena $[\![\theta_1]\!] \times \dots \times [\![\theta_n]\!] \Rightarrow [\![\theta]\!]$. The identity strategies are used to interpret free identifiers. Other elements of the syntax are interpreted by composition with designated strategies. Below we give plays defining some of them (as the least *saturated* strategies containing the plays). We use subscripts to indicate the subarena a move comes from.

	Arena	Generators
;	$[\![\text{com}]\!]_0 \times [\![\beta]\!]_1 \Rightarrow [\![\beta]\!]_2$	$q_2 \text{run}_0 \text{done}_0 q_1 a_1 a_2$
	$[\![\text{com}]\!]_0 \times [\![\text{com}]\!]_1 \Rightarrow [\![\text{com}]\!]_2$	$\text{run}_2 \text{run}_0 \text{run}_1 \text{done}_0 \text{done}_1 \text{done}_2$
:=	$[\![\text{var}]\!]_0 \times [\![\text{exp}]\!]_1 \Rightarrow [\![\text{com}]\!]_2$	$\text{run}_2 q_1 m_1 \text{write}(m)_0 \text{ok}_0 \text{done}_2$
!	$[\![\text{var}]\!]_0 \Rightarrow [\![\text{exp}]\!]_1$	$q_1 \text{read}_0 m_0 m_1$
grab	$[\![\text{sem}]\!]_0 \Rightarrow [\![\text{com}]\!]_1$	$\text{run}_1 \text{grab}_0 \text{ok}_0^g \text{done}_1$
release	$[\![\text{sem}]\!]_0 \Rightarrow [\![\text{com}]\!]_1$	$\text{run}_1 \text{release}_0 \text{ok}_0^r \text{done}_1$
newvar	$q_2 q_1 (\text{read}_0 0_0)^* (\sum_{i \in \mathbb{N}} (\text{write}(i)_0 \text{ok}_0 (\text{read}_0 i_0)^*))^* a_1 a_2$	
newsem	$q_2 q_1 (\text{grab}_0 \text{ok}_0^g \text{release}_0 \text{ok}_0^r)^* (\text{grab}_0 \text{ok}_0^g + \epsilon) a_1 a_2$	

The strategies for variable- and semaphore-binding are for playing in arenas $([\![\text{var}]\!]_0 \Rightarrow [\![\beta]\!]_1) \Rightarrow [\![\beta]\!]_2$ and $([\![\text{sem}]\!]_0 \Rightarrow [\![\beta]\!]_1) \Rightarrow [\![\beta]\!]_2$ respectively.

As shown in [4], the interpretation of PA sketched above yields a fully abstract model as detailed in Theorem 3.5. A play is called **complete** if it does not contain

unanswered questions. We write $\text{comp}(\sigma)$ to denote the set of non-empty *complete* plays of the strategy σ .

Theorem 3.5 [4] *For any PA-terms $\Gamma \vdash M_1 : \theta$ and $\Gamma \vdash M_2 : \theta$, $\Gamma \vdash M_1 \sqsubseteq_{\text{PA}} M_2$ if, and only if, $\text{comp}(\llbracket \Gamma \vdash M_1 \rrbracket) \subseteq \text{comp}(\llbracket \Gamma \vdash M_2 \rrbracket)$. Hence, $\Gamma \vdash M_1 \cong_{\text{PA}} M_2$ if, and only if, $\text{comp}(\llbracket \Gamma \vdash M_1 \rrbracket) = \text{comp}(\llbracket \Gamma \vdash M_2 \rrbracket)$.*

We are going to prove an analogous result for \mathbf{IA}^{\parallel} , though the preorder involved will be much more complicated.

4 Cloning

A shuffled sequence which is an interleaving of plays will be called a shuffled play. A shuffled play will be called *complete* if it is an interleaving of complete plays. In order to capture contextual approximation in \mathbf{IA}^{\parallel} , it turns out useful to introduce an auxiliary operation on complete shuffled plays. The operation will clone part of the sequence, namely, a selected question along with all the moves that it justifies.

Formally, let s be a complete shuffled play and let q be an occurrence of a question in s . Suppose m_1, \dots, m_k are all the moves hereditarily justified by q in s and, in particular, that m_k is the answer justified by q . For convenience we write m_0 for q , so that $s = s_0 m_0 s_1 m_1 \dots s_k m_k s_{k+1}$, where each s_i ($0 \leq i \leq k+1$) is a possibly empty *sequence* of moves. Let us now define another sequence s_q to be s in which each m_i ($0 \leq i \leq k$) is followed by its fresh copy m'_i , i.e.

$$s_q = s_0 m_0 m'_0 s_1 m_1 m'_1 \dots s_k m_k m'_k s_{k+1},$$

m_0 and m'_0 are justified by the same move (from s_0 , if any) and m'_i justifies m'_j ($i < j$) if, and only if, m_i justifies m_j . We shall call m'_0 and m'_k the **anchor points**. Intuitively, s_q can be thought of as s in which part of the play is being “shadowed”, as in a racing computation. Note that if s is a complete play and q is chosen to be the initial question, then the whole of s will be cloned and s_q will become a complete shuffled play.

Definition 4.1 Given two complete shuffled plays $s, t \in P_A$, we shall write $s \ll t$ provided s contains an occurrence of a question q such that $t = s_q$. If we want to stress that q is an X -question ($X \in \{O, P\}$), we write $s \ll_X t$. In what follows, we shall often consider the transitive closure of the above relations, which will be denoted by \ll^* , \ll_O^* and \ll_P^* respectively.

Example 4.2 (i) Consider the following two plays in $(\llbracket \mathbf{com} \rrbracket_3 \Rightarrow \llbracket \mathbf{com} \rrbracket_2) \times \llbracket \mathbf{exp} \rrbracket_1 \Rightarrow \llbracket \mathbf{com} \rrbracket_0$.

$$\begin{array}{cccccccc} s_1 = & r_0 & r_2 & \widehat{r_3} & q_1 & 0_1 & d_3 & d_2 & d_0 \\ & O & P & O & P & O & P & O & P \end{array} \quad \begin{array}{cccccccc} s_2 = & r_0 & r_2 & \widehat{r_3} & \widehat{r_3} & q_1 & 0_1 & d_3 & d_3 & d_2 & d_0 \\ & O & P & O & O & P & O & P & P & O & P \end{array}$$

We have omitted some pointers for the sake of clarity: in both plays r_0 justifies r_2, q_1, d_0 ; r_2 justifies d_2 ; r_0 justifies d_0 , and q_1 justifies 0_1 . Then $s_1 \ll_O s_2$.

(ii) Consider the following two plays in $\llbracket \mathbf{com} \rrbracket_1 \Rightarrow \llbracket \mathbf{com} \rrbracket_0$.

$$s_1 = r_0 \overset{\curvearrowright}{r_1} \overset{\curvearrowright}{d_1} d_0 \qquad s_2 = r_0 \overset{\curvearrowright}{r_1} \overset{\curvearrowright}{r_1} \overset{\curvearrowright}{d_1} \overset{\curvearrowright}{d_1} d_0$$

Note that $s_1 \ll_P s_2$.

Definition 4.3 Let $\sigma_1, \sigma_2 : A$. We define $\sigma_1 \leq \sigma_2$ to hold when for any $s_1 \in \text{comp}(\sigma_1)$ there exists $s_2 \in \text{comp}(\sigma_2)$ such that $s_1 \ll_P^* s_2$.

Example 4.4 $\llbracket x : \text{com} \vdash x : \text{com} \rrbracket \leq \llbracket x : \text{com} \vdash x \rrbracket \llbracket x : \text{com} \rrbracket$

\leq underpins our full abstraction result. The remainder of the paper will be devoted to its proof.

Theorem 4.5 (Full Abstraction) Let $\Gamma \vdash M_1, M_2 : \theta$ be IA^\parallel -terms. Then $\Gamma \vdash M_1 \sqsubseteq_{\text{IA}^\parallel} M_2$ if, and only if, $\llbracket \Gamma \vdash M_1 \rrbracket \leq \llbracket \Gamma \vdash M_2 \rrbracket$.

5 Definability

First we proceed to establish the left-to-right implication of Theorem 4.5, for which we need to prove a definability result. Recall from [4] that, for any complete play s , it is possible to construct a PA-term such that the corresponding strategy is the least saturated strategy containing s . This property no longer holds for IA^\parallel -terms (this will follow from the next section in which we identify a closure property of strategies corresponding to IA^\parallel -terms). Instead we shall prove a weakened result for IA^\parallel (Lemma 5.2).

Example 5.1 Let us write $[cond]$ for **if** $cond$ **then** **skip** **else** Ω_{com} . Consider the play $s = r_0 r_1 r_2 d_2 d_1 d_0$ from $\llbracket (\text{com}_2 \rightarrow \text{com}_1) \rightarrow \text{com}_0 \rrbracket$ and the term

$$\lambda f^{\text{com} \rightarrow \text{com}}. \text{newvar } X \text{ in newsem } S \text{ in } f(\text{grab}(S); X := 1); [!X = 1],$$

which is actually interpreted by the least saturated strategy containing s . When semaphores are no longer available, the “best” one can do to make sure that the assignment $X := 1$ is executed once is to protect it with the guard $[!X = 0]$ instead of $\text{grab}(S)$. However, this will not prevent multiple assignments from taking place if f runs several copies of its argument in parallel (so that each can pass the test $!X = 0$ before X is set to 1). Accordingly, the strategy corresponding to

$$\lambda f^{\text{com} \rightarrow \text{com}}. \text{newvar } X \text{ in } f([!X = 0]; X := 1); [!X = 1],$$

will contain, among others, the complete play $r_0 r_1 r_2 r_2 d_2 d_2 d_1 d_0$. In fact, the strategy contains all complete plays t such that $s \ll_O^* t$. This observation admits the following generalization.

Lemma 5.2 Suppose Θ is an IA^\parallel -type and $s \in \text{comp}(P_{\llbracket \Theta \rrbracket})$. Then there exists an IA^\parallel -term $\vdash M_s : \Theta$ such that

$$\text{comp}(\llbracket \vdash M_s \rrbracket) = \{ u \mid \exists t \in P_{\llbracket \Theta \rrbracket}. (s \ll_O^* t \text{ and } t \preceq u) \}.$$

Let us describe some of the ideas behind the construction of M_s . First of all, it is worth noting that, since saturated strategies are involved, s determines dependencies of P-moves on preceding O-moves. To enforce that order, we arrange for O-moves

to generate global side-effects ($G_i := 1$) so that P-moves can only take place if the side-effects corresponding to preceding O-moves occurred.

The example above shows that with shared memory alone we are unable to control the exact number of O-moves in complete plays. However, we can make sure that whenever copies of O-moves from the original play are played, they are globally synchronized. To this end, before the corresponding flag variable G_i is set to 1, we arrange for a test $[!G_i = 0]$. This creates a “window of opportunity” for the racing O-moves, into which they have to fit if a complete play is to be reached (late arrivals will fail the test and cause divergence).

Having synchronized racing on O-moves, we also need to make sure that the “races” are consistent with s . The global side effects are not enough for that purpose as they only signal that in *one* of the races the requisite moves have been made. To ensure consistency with s in cloned subplays (i.e. to ensure that all relevant moves from s are cloned) we introduce local flags L_i , each of which is set at the same time as G_i , except that there is a *local* test whether L_i has indeed been set. It suffices to use this mechanism for O-questions only, as the presence of O-answers follows from the fact that a complete play is to be reached in the end.

Example 5.3 Consider $\theta \equiv (((\mathbf{com}_4 \rightarrow \mathbf{com}_3) \rightarrow \mathbf{com}_2) \rightarrow \mathbf{com}_1) \rightarrow \mathbf{com}_0$ and the following play $s \in P_{[\theta]}$, in which we suppressed pointers from questions to answers.

$$\begin{array}{cccccccccccc} r_0 & \widehat{r_1} & \widehat{r_2} & \widehat{r_3} & \widehat{r_4} & d_4 & d_3 & d_2 & d_1 & d_0 \\ O & P & O & P & O & P & O & P & O & P \end{array}$$

The term M_s below satisfies Lemma 5.2. Note how the presence of L_4 ensures that in any complete play from $\llbracket \vdash M_s \rrbracket$ containing two occurrences of r_2 we must also have at least one occurrence of r_4 hereditarily justified by r_2 . G_4 alone would not suffice for this purpose.

$\lambda f.$ **newvar** $G_0, G_2, G_4, G_6, G_8, L_0$ **in**

$[!G_0 = 0]; G_0 := 1; L_0 := 1; [!G_0 = 1];$

newvar L_2 **in**

$f(\lambda g. [!G_2 = 0]; G_2 := 1; L_2 := 1; [\bigwedge_{j \in \{0,2\}} (!G_j = 1)]);$

newvar L_4 **in**

$g([!G_4 = 0]; G_4 := 1; L_4 := 1; [\bigwedge_{j \in \{0,2,4\}} (!G_j = 1)]);$

$[!L_4 = 1];$

$[!G_6 = 0]; G_6 := 1; [\bigwedge_{j \in \{0,2,4,6\}} (!G_j = 1)];$

$[!L_2 = 1];$

$[!G_8 = 0]; G_8 := 1; [\bigwedge_{j \in \{0,2,4,6,8\}} (!G_j = 1)]$

With the definability result in place, we obtain the following corollary.

Corollary 5.4 *For any \mathbf{IA}^{\parallel} -terms $\vdash M_1, M_2 : \theta$, if $\vdash M_1 \sqsubseteq_{\mathbf{IA}^{\parallel}} M_2 : \theta$ then $\llbracket \vdash M_1 : \theta \rrbracket \leq \llbracket \vdash M_2 : \theta \rrbracket$.*

6 Soundness

In this section we identify a technical property satisfied by strategies corresponding to \mathbf{IA}^{\parallel} -terms. In addition to helping us complete the proof of our full abstraction result, it provides us with a tool for checking whether a given strategy might originate from an \mathbf{IA}^{\parallel} -term.

Lemma 6.1 *Let $\Gamma \vdash M$ be an \mathbf{IA}^{\parallel} -term, $\sigma = \llbracket \Gamma \vdash M \rrbracket$ and $s \in \text{comp}(\sigma)$. Then, for any play t such that $s \ll_O^* t$, there exists $u \in \text{comp}(\sigma)$ such that $t \ll_P^* u$.*

Intuitively, the Lemma asserts that, for each successful interaction between the environment and the system, the environment can always trigger others, which closely follow (race) the original blueprint. Its logical structure resembles the conditions used to characterize **mkvar**-free computation in the game semantics literature [8,9].

Before discussing the proof, let us consider a number of examples.

Example 6.2 (i) Lemma 6.1 fails for the strategy σ used to interpret semaphore-binding, generated by plays of the form

$$q_2 q_1 (grab_0 ok_0^g release_0 ok_0^r)^* (grab_0 ok_0^g + \epsilon) a_1 a_2.$$

Observe that $s = q_2 q_1 grab_0 ok_0^g a_1 a_2 \in \sigma$ and consider $t = q_2 q_1 grab_0 grab_0 ok_0^g ok_0^g a_1 a_2$. Clearly $s \ll_O t$. However, note that $t \ll_P^* u$, where $u \in \text{comp}(\sigma)$, must imply $t = u$ (the only P-move that can possibly be taken to support $t \ll_P^+ u$ is q_1 , but plays in σ can only contain one occurrence of q_1). $t = u \in \text{comp}(\sigma)$ is impossible, though, because any play from σ that contains two occurrences of ok_0^g must contain at least one occurrence of $release_0$. Consequently, the “semaphore strategy” does not satisfy Lemma 6.1.

- (ii) The reasoning above does *not* apply to the strategy τ responsible for memory management. For instance, for $s = q_2 q_1 write(3) ok a_1 a_2$, $t = q_2 q_1 write(3) write(3) ok ok a_1 a_2$ we do have $t \in \text{comp}(\tau)$, because one of the defining plays is $q_2 q_1 write(3) ok write(3) ok a_1 a_2$.
- (iii) The identity strategy is easily seen to satisfy Lemma 6.1.
- (iv) All the other strategies corresponding to the syntax of \mathbf{IA}^{\parallel} satisfy the Lemma vacuously, because $s \in \text{comp}(\sigma)$ and $s \ll_O^* t$, where t is a play, imply $s = t$.

To prove the Lemma it suffices to show that the property involved is preserved by composition. The natural approach would be to try to apply the property to the two strategies alternately with the hope of deriving it for the composite. However, given the current formulation, this alternation might seemingly have no end! To recover, we shall make the property more precise by relating the operations witnessing $t \ll_P^* u$ to those fulfilling the same task for $s \ll_O^* t$. Intuitively, we want to express the fact that each of the clonings underlying $t \ll_P^* u$ is embedded into a cloning underpinning $s \ll_O^* t$. To make the intuition precise, let us assign a fresh colour to the two anchor points involved in each step of $s \ll_O^* t$ (the colours are to stay with the moves as additional moves are being added). Then we shall say that

$t \ll_P^* u$ *occurs within* $s \ll_O^* t$ iff for each pair of anchor points generated during the passage from t to u (according to $t \ll_P^* u$), both are between moves of the same colour.

An immediate consequence of the new requirement will be that the maximum distance (calculated in a way to be introduced) between anchor points involved in $s \ll_O^* t$ will be strictly larger than the maximum distance between anchor points generated by $t \ll_P^* u$ ². This is not necessarily the case for the obvious notion of distance (number of moves in-between), because \ll -steps add moves to plays.

Definition 6.3 Given a sequence of moves s , we define the *alternating length* of s to be the number of times the ownership of moves changes as we scan the sequence from left to right. The empty sequence is assumed to have alternating length 0.

For instance, $o_1o_2o_3$ is of (alternating) length 0, $o_1o_2p_2p_3$ has length 1 and $o_1p_1o_2p_3$ is of length 3. From now on, the distance between anchor points will be defined to be the *alternating length* of the segment between them (without the points). Given $s_1 \ll_X^* s_2$ we shall say that the associated **weight** is the largest of the distances between anchor points involved in the transitions from s_1 to s_2 . Note that if $s \ll t$ then s and t have the same alternating length. Because of that, if $t \ll_P^* u$ occurs within $s \ll_O^* t$, the weight of $t \ll_P^* u$ must be strictly smaller than that of $s \ll_O^* t$.

Another consequence of “occurring within”, crucial for establishing compositionality, is the fact that during composition of σ with τ , due to the embeddings, local decreases in weight effected by σ imply that the corresponding weight calculated for σ^* also decreases. Moreover, the decreases caused by σ and τ can be meaningfully combined. As a consequence, we can show a strengthened version of Lemma 6.1.

Lemma 6.4 Let $\Gamma \vdash M$ be an \mathbf{IA}^{\parallel} -term, $\sigma = \llbracket \Gamma \vdash M \rrbracket$ and $s \in \text{comp}(\sigma)$. Then, for any play t such that $s \ll_O^* t$, there exists $u \in \text{comp}(\sigma)$ such that $t \ll_P^* u$, and $t \ll_P^* u$ occurs within $s \ll_O^* t$.

Example 6.5 The closure property spelt out in Lemma 6.4 shows that the problems identified in Example 5.1 are unavoidable: there can be no \mathbf{IA}^{\parallel} -term $\vdash M : ((\mathbf{com}_2 \rightarrow \mathbf{com}_1) \rightarrow \mathbf{com}_0)$ such that $\text{comp}(\llbracket \vdash M \rrbracket) = \{r_0r_1r_2d_2d_1d_0\}$. In contrast, the PA-term

$$\lambda f^{\mathbf{com} \rightarrow \mathbf{com}}. \text{newvar } X \text{ in newsem } S \text{ in } f(\text{grab}(S); X := 1); [!X = 1]$$

does satisfy the equation. Consequently, in PA, semaphores (in fact, even a single occurrence of **grab**) cannot be replaced with shared memory up to observational equivalence.

Example 6.6 The test-and-set instruction (**test-set**(X)) sets the value of the given variable to 1 and returns the old value as a single atomic (non-interruptible) operation. Observe that, if we added **test-set**(X) to \mathbf{IA}^{\parallel} , we could replace

$$\lambda f^{\mathbf{com} \rightarrow \mathbf{com}}. \text{newvar } X \text{ in newsem } S \text{ in } f(\text{grab}(S); X := 1); [!X = 1]$$

² Abusing the notation somewhat, here we regard $s \ll_O^* t$ as shorthand for a concrete sequence demonstrating that $s \ll_O^* t$.

with

$$\lambda f^{\mathbf{com} \rightarrow \mathbf{com}}. \mathbf{newvar} \, X \, \mathbf{in} \, f([\mathbf{test-set}(X) = 0]); [!X = 1].$$

Since the definability argument for PA [4] only relies on “grabs” of this kind, it carries over to $\mathbf{IA}^{\parallel} + \mathbf{test-set}$. Consequently, $\mathbf{IA}^{\parallel} + \mathbf{test-set}$ has the same discriminating power as PA.

Using Lemma 6.4 we can eventually complete the proof of Theorem 4.5 by showing :

Corollary 6.7 *For any \mathbf{IA}^{\parallel} -terms $\vdash M_1 : \theta$ and $\vdash M_2 : \theta$, if $\llbracket \vdash M_1 : \theta \rrbracket \leq \llbracket \vdash M_2 : \theta \rrbracket$ then $\vdash M_1 \sqsubseteq_{\mathbf{IA}^{\parallel}} M_2 : \theta$.*

7 Conclusion

We have constructed an inequationally fully abstract model of \mathbf{IA}^{\parallel} inside an existing model of PA and given an explicit characterization of contextual approximation in PA in terms of a preorder on complete plays. We have also identified a closure property that all \mathbf{IA}^{\parallel} -terms satisfy and some PA-terms do not. Consequently, we can conclude that semaphores cannot be programmed in \mathbf{IA}^{\parallel} if the translation is to preserve observational equivalence (of PA-terms). So, why do the solutions to the mutual exclusion problem not apply?

The reason is that semaphores offer a *uniform* solution to the mutual exclusion problem. Whenever different processes intend to use a critical section, they can run identical entry and exit protocols (**grab**(S) and **release**(s) respectively). In contrast, existing solutions based on shared memory are not uniform, even though they are often “symmetric”, in that the code run by each process depends only on its identifier. For instance, in Peterson’s algorithm the codes for the two processes are the same up to the permutation that swaps 1 and 2. Such solutions will not help us to mimic the effect of **grab**(S) in, say, $f(\mathbf{grab}(S))$, because f can also make its argument run in parallel with itself, a scenario which does not arise in the framework of cooperating sequential processes.

Furthermore, our results demonstrate that PA is not a conservative extension of \mathbf{IA}^{\parallel} with respect to observational equivalence (and hence also observational approximation). Here are the simplest instances of that failure, now easily verifiable, thanks to Theorems 3.5 and 4.5.

- (i) $x : \mathbf{com} \vdash x \sqsubseteq_{\mathbf{IA}^{\parallel}} (x||x) : \mathbf{com}$
 $x : \mathbf{com} \vdash x \not\sqsubseteq_{\mathbf{PA}} (x||x) : \mathbf{com}$
- (ii) $x : \mathbf{com} \vdash (x \mathbf{or} (x||x)) \cong_{\mathbf{IA}^{\parallel}} (x||x) : \mathbf{com}$
 $x : \mathbf{com} \vdash (x \mathbf{or} (x||x)) \not\cong_{\mathbf{PA}} (x||x) : \mathbf{com}$

Acknowledgement

I am grateful to Samson Abramsky for discussions on the topic of this paper.

References

- [1] Abramsky, S. and G. McCusker, *Linearity, sharing and state: a fully abstract game semantics for Idealized Algol with active expressions*, in: P. W. O'Hearn and R. D. Tennent, editors, *Algol-like languages*, Birkhäuser, 1997 pp. 297–329.
- [2] Brookes, S. D., *The essence of Parallel Algol*, *Information and Computation* **179** (2002), pp. 118–149.
- [3] Dijkstra, E. W., *Cooperating sequential processes*, in: F. Genuys, editor, *Programming Languages: NATO Advanced Study Institute*, Academic Press, 1968 pp. 43–112.
- [4] Ghica, D. R. and A. S. Murawski, *Angelic semantics of fine-grained concurrency*, *Annals of Pure and Applied Logic* **151(2-3)** (2008), pp. 89–114.
- [5] Jifeng, H., M. B. Josephs and C. A. R. Hoare, *A theory of synchrony and asynchrony*, in: *Programming Concepts and Methods*, Elsevier, 1990 pp. 459–473.
- [6] Laird, J., *A game semantics of Idealized CSP*, in: *Proceedings of MFPS'01*, Elsevier, 2001 pp. 1–26, ENTCS, Vol. 45.
- [7] Lamport, L., *A new solution of Dijkstra's concurrent programming problem*, *Communications of the ACM* **17** (1974), pp. 453–455.
- [8] McCusker, G., *On the semantics of Idealized Algol without the bad-variable constructor.*, in: *Proceedings of MFPS*, *Electronic Notes in Theoretical Computer Science* **83** (2003).
- [9] Murawski, A. S., *Bad variables under control*, in: *Proceedings of CSL*, *Lecture Notes in Computer Science* **4646**, Springer, 2007 pp. 558–572.
- [10] Peterson, G. L., *Myths about the mutual exclusion problem*, *Information Processing Letters* **12** (1981), pp. 115–116.
- [11] Reynolds, J. C., *The essence of Algol*, in: J. W. de Bakker and J. van Vliet, editors, *Algorithmic Languages*, North Holland, 1981 pp. 345–372.
- [12] Röckl, C. and D. Sangiorgi, *A pi-calculus process semantics of Concurrent Idealised Algol*, in: *Proceedings of FoSSaCS*, *Lecture Notes in Computer Science* **1578** (1999), pp. 306–321.
- [13] Udding, J. T., *A formal model for defining and classifying delay-insensitive circuits and systems*, *Distributed Computing* **1(4)** (1986), pp. 197–204.